

Full Strategic Roadmap Report

GPU TARGET

H100

BREAKEVEN DAY

364 days

OWNERSHIP TOTAL COST

\$28,750

Infrastructure Breakeven Analysis

Purchase price	\$25,000
Hourly rental rate	\$6.88 / hr (On-demand)
Monthly rental equivalent	\$2,064
Electricity/cooling buffer	\$3,750
Term rental cost	\$37,152

API vs Self-Hosted Comparison

API model	GPT-5
Reasoning overhead	20% output-token overhead applied
Daily requests	60,000
Monthly API burn	\$16,515
Project API cost	\$297,270
Monthly token volume	5,130,000,000

Source Model Metadata

GPU category	Datacenter
GPU VRAM	80 GB
Hardware cost basis	\$25,000
On-demand rental	\$6.88 / hr
Spot rental	\$2.19 / hr

API provider	OpenAI
API tier	Frontier Reasoning
Input price	\$1.25 / 1M tokens
Output price	\$10.00 / 1M tokens
Prompt cache assumption	50% input-token discount

SCALABILITY VERDICT

High token volume favors dedicated infrastructure once utilization is stable.