

Full Strategic Roadmap Report

GPU TARGET

H100

BREAKEVEN DAY

834 days

OWNERSHIP TOTAL COST

\$28,750

Infrastructure Breakeven Analysis

Purchase price	\$25,000
Hourly rental rate	\$3.00 / hr (On-demand)
Monthly rental equivalent	\$900.00
Electricity/cooling buffer	\$3,750
Term rental cost	\$16,200

API vs Self-Hosted Comparison

API model	GPT-5
Reasoning overhead	20% output-token overhead applied
Daily requests	18,000
Monthly API burn	\$7,803
Project API cost	\$140,454
Monthly token volume	1,539,000,000

Source Model Metadata

GPU category	Datacenter
GPU VRAM	80 GB
Hardware cost basis	\$25,000
On-demand rental	\$3.00 / hr
Spot rental	\$1.80 / hr

API provider	OpenAI
API tier	Frontier Reasoning
Input price	\$2.50 / 1M tokens
Output price	\$15.00 / 1M tokens
Prompt cache assumption	50% input-token discount

SCALABILITY VERDICT

Hybrid deployment is likely optimal: reserve infrastructure for steady workloads and APIs for burst.